

Open VSwitch

2022-04-20

王定山

PPT: <https://shorturl.at/psuTV>

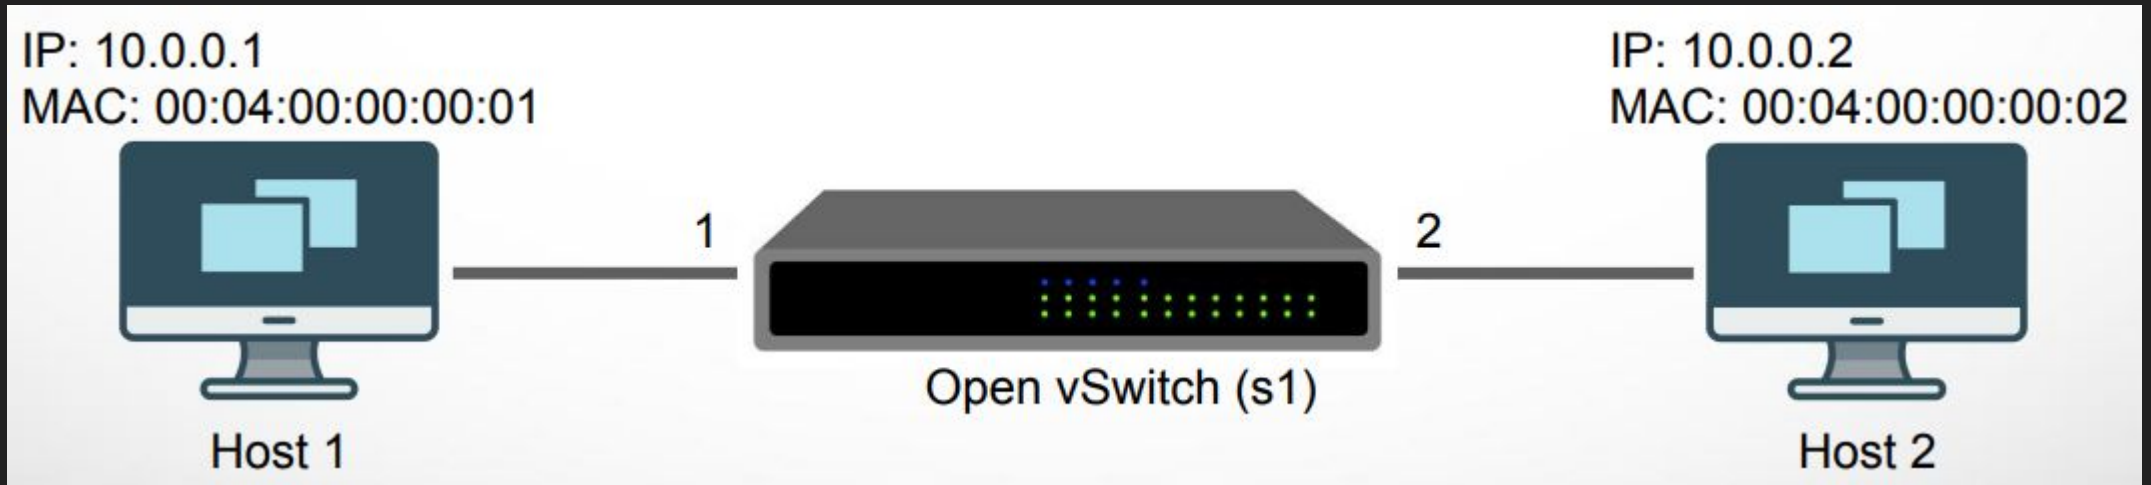
Meter

- A switch element that can measure and control the rate of packets.
- The meter trigger a meter band if the packet rate or byte rate passing through the meter exceed a predefined threshold.
- If the meter band drops the packet, it is called a **Rate Limiter**.



Meter

- Add meter entry (set maximum rate with 30 kbps)
`sudo ovs-ofctl add-meter -O OpenFlow13 s1 meter=1,kbps,band=type=drop,rate=30000`
- Add a flow entries with binding meter
`sudo ovs-ofctl add-flow -O OpenFlow13 s1 in_port=1,actions=meter:1,output:2`
`sudo ovs-ofctl add-flow -O OpenFlow13 s1 in_port=2,actions=meter:1,output:1`



Meter

```
mountain@openflow: ~/MininetTopology 63x4
mininet> iperf h1 h2
*** Iperf: testing TCP bandwidth between h1 and h2
*** Results: ['40.9 Mbits/sec', '41.4 Mbits/sec']
mininet>


mountain@openflow: ~/MininetTopology 63x10
mountain@openflow:~/MininetTopology$ bash ./answer-TestbedA.sh
OFPST_METER_CONFIG reply (OF1.3) (xid=0x2):
meter=1 kbps bands=
type=drop rate=30000
  cookie=0x0, duration=0.068s, table=0, n_packets=0, n_bytes=0,
in_port="s1-eth1" actions=meter:1,output:"s1-eth2"
  cookie=0x0, duration=0.040s, table=0, n_packets=0, n_bytes=0,
in_port="s1-eth2" actions=meter:1,output:"s1-eth1"
mountain@openflow:~/MininetTopology$
```

```
mountain@openflow: ~/MininetTopology 65x4
mininet> iperf h1 h2
*** Iperf: testing TCP bandwidth between h1 and h2
*** Results: ['95.7 Mbits/sec', '96.9 Mbits/sec']
mininet>

mountain@openflow: ~/MininetTopology 65x10
mountain@openflow:~/MininetTopology$ sudo ovs-ofctl dump-flows s1
  cookie=0x0, duration=13.602s, table=0, n_packets=0, n_bytes=0, in
n_port="s1-eth2" actions=output:"s1-eth1"
  cookie=0x0, duration=9.997s, table=0, n_packets=0, n_bytes=0, in
_port="s1-eth1" actions=output:"s1-eth2"
mountain@openflow:~/MininetTopology$
```

Groups

- Groups represent sets of actions for flooding, as well as more complex forwarding semantics (e.g. multipath, fast reroute, and link aggregation).
- As a general layer of indirection, groups also enable multiple flow entries to forward to a single identifier (e.g. IP forwarding to a common next hop).
- This abstraction allows common output actions across flow entries to be changed efficiently.



Groups


- Add group entry:

```
sudo ovs-ofctl add-group s1
```

```
group_id=1,type=all,bucket=mod_dl_src=00:00:00:11:11:11,mod_dl_dst=00:00:00:22:22:22,output:2,bucket=mod_dl_src=00:00:00:11:11:11,mod_dl_dst=00:00:00:22:22:22,output:3,bucket=mod_dl_src=00:00:00:11:11:11,mod_dl_dst=00:00:00:22:22:22,output:4
```

- Add flow entry with using group action:

```
sudo ovs-ofctl add-flow s1 in_port=1,actions=group:1
```



Groups

```
mountain@openflow:~/MininetTopology$ sudo ovs-ofctl add-group s1 group_id=1,type=all,bucket=mod_dl_src=00:00:00:11:11:11,mod_dl_dst=00:00:00:22:22:22,output:2,bucket=mod_dl_src=00:00:00:11:11:11,mod_dl_dst=00:00:00:22:22:22,output:3,bucket=mod_dl_src=00:00:00:11:11:11,mod_dl_dst=00:00:00:22:22:22,output:4
mountain@openflow:~/MininetTopology$ sudo ovs-ofctl add-flow s1 in_port=1,actions=group:1
mountain@openflow:~/MininetTopology$ sudo ovs-ofctl dump-groups s1
NXST_GROUP_DESC reply (xid=0x2):
  group_id=1,type=all,bucket=bucket_id:0,actions=mod_dl_src:00:00:00:11:11:11,mod_dl_dst:00:00:00:22:22:22,output:"s1-eth2",bucket=bucket_id:1,actions=mod_dl_src:00:00:00:11:11:11,mod_dl_dst:00:00:00:22:22:22,output:"s1-eth3",bucket=bucket_id:2,actions=mod_dl_src:00:00:00:11:11:11,mod_dl_dst:00:00:00:22:22:22,output:"s1-eth4"
mountain@openflow:~/MininetTopology$ sudo ovs-ofctl dump-flows s1
cookie=0x0, duration=21.017s, table=0, n_packets=1, n_bytes=70, in_port="s1-eth1" actions=group:1
```

```
mountain@openflow:~/MininetTopology$ sudo tcpdump -i s1-eth1 -e
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on s1-eth1, link-type EN10MB (Ethernet), capture size 262144 bytes
15:08:06.056193 00:04:00:00:00:01 (oui Unknown) > Broadcast, ethertype ARP (0x0806), length 42: Request who-has 10.0.0.1 tell 10.0.0.1, length 28
15:08:07.076258 00:04:00:00:00:01 (oui Unknown) > Broadcast, ethertype ARP (0x0806), length 42: Request who-has 10.0.0.1 tell 10.0.0.1, length 28
```

Groups

```
mountain@openflow:~/MininetTopology$ sudo tcpdump -i s1-eth2 -e
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on s1-eth2, link-type EN10MB (Ethernet), capture size 262144 bytes
15:07:58.884461 00:00:00:11:11:11 (oui Ethernet) > 00:00:00:22:22:22 (oui Ethernet), ethertype ARP (0x0806), length 42:
Request who-has 10.0.0.2 tell 10.0.0.1, length 28
```

```
mountain@openflow:~/MininetTopology$ sudo tcpdump -i s1-eth3 -e
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on s1-eth3, link-type EN10MB (Ethernet), capture size 262144 bytes
15:07:49.668456 00:00:00:11:11:11 (oui Ethernet) > 00:00:00:22:22:22 (oui Ethernet), ethertype ARP (0x0806), length 42:
Request who-has 10.0.0.2 tell 10.0.0.1, length 28
```

```
mountain@openflow:~/MininetTopology$ sudo tcpdump -i s1-eth4 -e
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on s1-eth4, link-type EN10MB (Ethernet), capture size 262144 bytes
15:08:31.652308 00:00:00:11:11:11 (oui Ethernet) > 00:00:00:22:22:22 (oui Ethernet), ethertype ARP (0x0806), length 42:
Request who-has 10.0.0.2 tell 10.0.0.1, length 28
```



Experiment: IPv4 Multicasting

- Goal:
 - 編寫交換機S1,2,3路由規則, 讓封包送到相同的Multicast group使用者。
- Multicast groups:
 - Group A: 10.0.0.0/24
 - Group B: 10.0.1.0/24
- 測試
 - 使用指令sudo python3 ./Lab2.py啟動Mininet測試環境
 - 使用xterm A1 B1 打開A1, B1之終端機
 - 於A1/B1終端機執TestScript_Lab1.py
 - A1: sudo python3 ./TestScript_Lab1.py 224.0.0.1
 - B1: sudo python3 ./TestScript_Lab1.py 224.0.0.2
 - 把記錄檔的內容擷圖表示即可。




Experiment: IPv4 Multicasting

- Topology




Homework - IPv4 Routing with Meter Rate limiter



- 使用OpenVSwitch 指令對以下三條連線加入路由規則，並使用Meter 限定流量。
- 使用sudo ovs-ofctl dump-flows <bridge_id>把每一台交換機上的的規則印出並擷圖。
- 要求：使用IPv4 位址及UDP Port ID進行路由 (不需要寫ARP)
- 路線及使用的UDP Port
 - 紅色：UDP Port ID 50001, meter id = 1, drop, max bandwidth: 50Mbps
 - 藍色：UDP Port ID 50002, meter id = 2, drop, max bandwidth: 20Mbps
 - 綠色：UDP Port ID 50003, meter id = 3, drop, max bandwidth: 30Mbps

Homework - IPv4 Routing

1. 啟用環境: `sudo python3 ./Homework.py`
2. 編寫路由規則 (Flow entries)
3. 於Mininet 使用xterm h1 打開Host 之終端機
4. 在h1的終端機上執行指令 `bash ./run_test.sh` 進行測試, 程式過程及結果會寫在log上。
5. 測試後把50001-s.log, 50002-s.log, 50003-s.log之檔案內容擷圖。



Hints:

編寫規則時要比對封包上ip_protocol 後才能比對UDP封包內容 (PPT P.30)